


Logistic Regression: Multiclass Classification

Problem:

Hand written recognition is the problem we are trying to solve. We will use a training set with a lot of hand digit characters and then we build a model using logistic regression.

Identify hand written digits recognition



for more information check this site (https://scikit-learn.org/stable/auto_examples/datasets/plot_digits_last_image.html).

```
In [2]: %matplotlib inline
import matplotlib.pyplot as plt

from sklearn.datasets import load_digits
```

```
In [3]: # To load my training set
digits = load_digits()
```

```
In [5]: # Explore what this training set contains?
dir(digits)
```

```
Out[5]: ['DESCR', 'data', 'feature_names', 'frame', 'images', 'target', 'target_names']
```

As you can see from the site above there are **1797** 8*8 samples.

```
In [8]: #print the first sample
digits.data[0]
```


```
Out[8]: array([ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.,  0.,  0., 13., 15., 10.,
                15.,  5.,  0.,  0.,  3., 15.,  2.,  0., 11.,  8.,  0.,  0.,  4.,
                12.,  0.,  0.,  8.,  8.,  0.,  0.,  5.,  8.,  0.,  0.,  9.,  8.,
                 0.,  0.,  4., 11.,  0.,  1., 12.,  7.,  0.,  0.,  2., 14.,  5.,
                10., 12.,  0.,  0.,  0.,  0.,  6., 13., 10.,  0.,  0.,  0.]
```

It's an array as such it is an 8*8 image but the image is represented as a 1 dimensional array. If you want to see this particular element you can use matplotlib.

```
In [11]: # Show an actual image for numeric data
plt.gray()
plt.matshow(digits.images[0])
```




```
Out[11]: <matplotlib.image.AxesImage at 0xa7647c0>
```


<Figure size 432x288 with 0 Axes>



```
In [12]: #Print the first 5 samples
plt.gray()
for i in range(5):
    plt.matshow(digits.images[i])
```

<Figure size 432x288 with 0 Axes>





```
In [14]: dir(digits)
```

```
Out[14]: ['DESCR', 'data', 'feature_names', 'frame', 'images', 'target', 'target_names']
```

```
In [15]: digits.target[0:5]
```

```
Out[15]: array([0, 1, 2, 3, 4])
```

We can use **data** and **target** to train our model.

```
In [17]: from sklearn.model_selection import train_test_split
```

```
In [18]: X_train, X_test, y_train, y_test = train_test_split(digits.data, digits.target, test_size=0.2)
```

```
In [19]: len(X_train)
```

```
Out[19]: 1437
```

$1437 = 0.8 \times 1797$

Create logistic regression model

```
In [23]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
```

```
In [24]: model.fit(X_train, y_train)
```

C:\Users\Ehsan\AppData\Roaming\Python\Python38\site-packages\sklearn\linear_model_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[24]: LogisticRegression()
```

Measure accuracy of our model


```
In [25]: model.score(X_test, y_test)
```

```
Out[25]: 0.9694444444444444
```

Use the model and predict [67]

```
In [27]: plt.matshow(digits.images[67])
```

```
Out[27]: <matplotlib.image.AxesImage at 0xa8cd6d0>
```



```
In [28]: digits.target[67]
```

```
Out[28]: 6
```

```
In [30]: model.predict(digits.data[[67]])
```

```
Out[30]: array([6])
```

Our model works fine.

```
In [31]: model.predict(digits.data[0:5])
```

```
Out[31]: array([0, 1, 2, 3, 4])
```

I want to know where my model fails in accuracy.

Confusion Matrix

```
In [32]: y_predicted = model.predict(X_test)
```


```
In [33]: from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, y_predicted)  
cm
```

```
Out[33]: array([[40,  0,  0,  0,  0,  1,  0,  0,  0,  0],  
               [ 0, 39,  0,  1,  1,  0,  0,  1,  0,  0],  
               [ 0,  0, 33,  0,  0,  0,  0,  0,  0,  0],  
               [ 0,  0,  0, 28,  0,  0,  0,  0,  0,  1],  
               [ 0,  0,  0,  0, 39,  0,  0,  0,  0,  0],  
               [ 0,  1,  0,  0,  0, 30,  0,  0,  0,  1],  
               [ 0,  0,  0,  0,  0,  1, 31,  0,  0,  0],  
               [ 0,  0,  0,  0,  0,  0,  0, 38,  0,  0],  
               [ 0,  1,  0,  0,  0,  0,  0,  1, 37,  0],  
               [ 0,  0,  0,  0,  0,  0,  0,  0,  1, 34]], dtype=int64)
```

This is better visualize in matplotlib and seaborn.

```
In [34]: import seaborn as sn
plt.figure(figsize = (10,7))
sn.heatmap(cm, annot=True)
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

Out[34]: Text(69.0, 0.5, 'Truth')



whenever the numbers other than main diagonal is not 0 means your model is inaccurate. For instance, the main diagonal is work fine 39 times the actual value and predicted one is 1. Number 5 is one time 1 and model inaccurate.


Exercise

Use `sklearn.datasets iris` flower dataset to train your model using logistic regression. You need to figure out accuracy of your model and use that to predict different samples in your test dataset. In iris dataset there are 150 samples containing following features,

1. Sepal Length
2. Sepal Width
3. Petal Length
4. Petal Width

Using above 4 features you will classify a flower in one of the three categories,

1. Setosa
2. Versicolour
3. Virginica



Tip: Use `from sklearn.datasets import load_iris`

[click here for more information for exercise \(https://scikit-learn.org/stable/modules/generated/sklearn.datasets.load_iris.html\)](https://scikit-learn.org/stable/modules/generated/sklearn.datasets.load_iris.html)

[Also check this site \(https://scikit-learn.org/stable/auto_examples/datasets/plot_iris_dataset.html\)](https://scikit-learn.org/stable/auto_examples/datasets/plot_iris_dataset.html)

Solution:

Import dependencies

```
In [1]: %matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
```

```
In [2]: #Load the data set
data = sns.load_dataset("iris")
data.head()
```

Out[2]:

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

```
In [3]: # Prepare the training set
# X = feature values, all of the columns except the last column
X = data.iloc[:, :-1]
X
```

Out[3]:

	sepal_length	sepal_width	petal_length	petal_width
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2
...
145	6.7	3.0	5.2	2.3
146	6.3	2.5	5.0	1.9
147	6.5	3.0	5.2	2.0
148	6.2	3.4	5.4	2.3
149	5.9	3.0	5.1	1.8

150 rows × 4 columns

```
In [4]: y=data.iloc[:,-1]
y
```

```
Out[4]: 0      setosa
1      setosa
2      setosa
3      setosa
4      setosa
...
145    virginica
146    virginica
147    virginica
148    virginica
149    virginica
Name: species, Length: 150, dtype: object
```

```
In [8]: # Plot the relation of each feature with each species
plt.xlabel('Features')
plt.ylabel('Species')


pltX=data.loc[:, 'sepal_length']
pltY=data.loc[:, 'species']
plt.scatter(pltX,pltY,color='b',label='sepal_length')

pltX=data.loc[:, 'sepal_width']
pltY=data.loc[:, 'species']
plt.scatter(pltX,pltY,color='g',label='sepal_width')

pltX=data.loc[:, 'petal_length']
pltY=data.loc[:, 'species']
plt.scatter(pltX,pltY,color='r',label='petal_length')

pltX=data.loc[:, 'petal_width']
pltY=data.loc[:, 'species']
plt.scatter(pltX,pltY,color='black',label='petal_width')

plt.legend(loc=4,prop={'size':8})
plt.show()
```



```
In [9]: # Split the data into 80% training and 20% testing
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)
```

```
In [11]: # Train the model
model=LogisticRegression()
model.fit(X_train,y_train)
```

C:\Users\Ehsan\AppData\Roaming\Python\Python38\site-packages\sklearn\linear_model_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

Out[11]: LogisticRegression()

```
In [13]: # Test the model
y_predict = model.predict(X_test)
y_predict
```

Out[13]: array(['versicolor', 'setosa', 'virginica', 'versicolor', 'versicolor',
 'setosa', 'versicolor', 'virginica', 'versicolor', 'versicolor',
 'virginica', 'setosa', 'setosa', 'setosa', 'setosa', 'versicolor',
 'virginica', 'versicolor', 'versicolor', 'virginica', 'setosa',
 'virginica', 'setosa', 'virginica', 'virginica', 'virginica',
 'virginica', 'virginica', 'setosa', 'setosa'], dtype=object)

```
In [14]: y_test
```

```
Out[14]: 73    versicolor
         18      setosa
        118    virginica
         78    versicolor
         76    versicolor
         31      setosa
         64    versicolor
        141    virginica
         68    versicolor
         82    versicolor
        110    virginica
         12      setosa
         36      setosa
          9      setosa
         19      setosa
         56    versicolor
        104    virginica
         69    versicolor
         55    versicolor
        132    virginica
         29      setosa
        127    virginica
         26      setosa
        128    virginica
        131    virginica
        145    virginica
        108    virginica
        143    virginica
         45      setosa
         30      setosa
        Name: species, dtype: object
```

```
In [15]: # Check precision, recall, f1-score
print(classification_report(y_test,y_predict))
```

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	1.00	1.00	1.00	9
virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
In [16]: # The accuracy
print(accuracy_score(y_test,y_predict))
```

```
1.0
```

